

Práctica AirBnB - Bloque I - Aprendizaje Automático

Jorge Marcos Puñal

Link al proyecto: <https://github.com/JotaEmme/AirBnB-Palma-study/blob/main/Airbnb.ipynb>

```
In [1]:  
# Librerías de datos  
import pandas as pd  
import numpy as np  
  
# Librerías gráficas  
import seaborn as sns  
from matplotlib import pyplot as plt  
  
# Librerías matemáticas  
import math  
import statistics  
  
# Librerías para el modelo de ML  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.tree import DecisionTreeRegressor  
from xgboost import XGBRegressor, to_graphviz  
from sklearn.feature_selection import SelectKBest, f_regression  
  
# Accuracy  
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error  
from sklearn.metrics import balanced_accuracy_score, roc_auc_score, make_scorer  
from sklearn.model_selection import GridSearchCV # cross validation
```

```
In [2]:  
# Carga del archivo  
raw_data = pd.read_csv("airbnb.csv", header = 0, index_col = 0)  
data = raw_data.copy()
```

Resumen del dataset

```
In [3]:  
data.head(2)
```

Out[3]:

	listing_url	scrape_id	last_scraped	name	des
id					
11547	https://www.airbnb.com/rooms/11547	20200919153121	2020-09-21	My home at the beach	pe
100831	https://www.airbnb.com/rooms/100831	20200919153121	2020-09-21	HOUSE IN MALLORCA - WiFi(ET-3045)	sp

2 rows × 73 columns

In [4]:

```
print("Dataset tiene {} filas y {} columnas.".format(*data.shape))
print("Column Information:", data.info())
```

```
Dataset tiene 17608 filas y 73 columnas.
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17608 entries, 11547 to 45499210
Data columns (total 73 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   listing_url      17608 non-null   object 
 1   scrape_id        17608 non-null   int64  
 2   last_scraped     17608 non-null   object 
 3   name             17607 non-null   object 
 4   description      17393 non-null   object 
 5   neighborhood_overview 8213 non-null   object 
 6   picture_url      17608 non-null   object 
 7   host_id          17608 non-null   int64  
 8   host_url          17608 non-null   object 
 9   host_name         17606 non-null   object 
 10  host_since        17606 non-null   object 
 11  host_location     17572 non-null   object 
 12  host_about        11696 non-null   object 
 13  host_response_time 15862 non-null   object 
 14  host_response_rate 15862 non-null   object 
 15  host_acceptance_rate 16098 non-null   object 
 16  host_is_superhost 17606 non-null   object 
 17  host_thumbnail_url 17606 non-null   object 
 18  host_picture_url  17606 non-null   object 
 19  host_neighbourhood 364 non-null    object 
 20  host_listings_count 17606 non-null   float64
 21  host_total_listings_count 17606 non-null   float64
 22  host_verifications 17608 non-null   object 
 23  host_has_profile_pic 17606 non-null   object 
 24  host_identity_verified 17606 non-null   object 
 25  neighbourhood      8213 non-null   object 
 26  neighbourhood_cleansed 17608 non-null   object 
 27  neighbourhood_group_cleansed 0 non-null    float64
 28  latitude           17608 non-null   float64
 29  longitude          17608 non-null   float64
 30  property_type      17608 non-null   object 
 31  room_type          17608 non-null   object 
 32  accommodates       17608 non-null   int64
```

```
33 bathrooms 0 non-null float64
34 bathrooms_text 17600 non-null object
35 bedrooms 17333 non-null float64
36 beds 17511 non-null float64
37 amenities 17608 non-null object
38 price 17608 non-null object
39 minimum_nights 17608 non-null int64
40 maximum_nights 17608 non-null int64
41 minimum_minimum_nights 17608 non-null int64
42 maximum_minimum_nights 17608 non-null int64
43 minimum_maximum_nights 17608 non-null int64
44 maximum_maximum_nights 17608 non-null int64
45 minimum_nights_avg_ntm 17608 non-null float64
46 maximum_nights_avg_ntm 17608 non-null float64
47 calendar_updated 0 non-null float64
48 has_availability 17608 non-null object
49 availability_30 17608 non-null int64
50 availability_60 17608 non-null int64
51 availability_90 17608 non-null int64
52 availability_365 17608 non-null int64
53 calendar_last_scraped 17608 non-null object
54 number_of_reviews 17608 non-null int64
55 number_of_reviews_ltm 17608 non-null int64
56 number_of_reviews_l30d 17608 non-null int64
57 first_review 11173 non-null object
58 last_review 11173 non-null object
59 review_scores_rating 10957 non-null float64
60 review_scores_accuracy 10951 non-null float64
61 review_scores_cleanliness 10953 non-null float64
62 review_scores_checkin 10949 non-null float64
63 review_scores_communication 10951 non-null float64
64 review_scores_location 10950 non-null float64
65 review_scores_value 10949 non-null float64
66 license 11431 non-null object
67 instant_bookable 17608 non-null object
68 calculated_host_listings_count 17608 non-null int64
69 calculated_host_listings_count_entire_homes 17608 non-null int64
70 calculated_host_listings_count_private_rooms 17608 non-null int64
71 calculated_host_listings_count_shared_rooms 17608 non-null int64
72 reviews_per_month 11173 non-null float64
dtypes: float64(19), int64(20), object(34)
memory usage: 9.9+ MB
Column Information: None
```

In [5]:

```
data.describe()
```

Out[5]:

	scrape_id	host_id	host_listings_count	host_total_listings_count	neighbo
count	1.760800e+04	1.760800e+04	17606.000000	17606.000000	
mean	2.020092e+13	1.012201e+08	130.200784	130.200784	
std	5.093895e+00	9.453596e+07	246.651175	246.651175	
min	2.020092e+13	4.294200e+04	0.000000	0.000000	
25%	2.020092e+13	2.116003e+07	2.000000	2.000000	
50%	2.020092e+13	8.063674e+07	15.000000	15.000000	
75%	2.020092e+13	1.516679e+08	111.000000	111.000000	
max	2.020092e+13	3.679802e+08	1136.000000	1136.000000	

8 rows × 39 columns

Entendimiento y limpieza del dataset

In [6]:

data.columns

```
Out[6]: Index(['listing_url', 'scrape_id', 'last_scraped', 'name', 'description',
       'neighborhood_overview', 'picture_url', 'host_id', 'host_url',
       'host_name', 'host_since', 'host_location', 'host_about',
       'host_response_time', 'host_response_rate', 'host_acceptance_rate',
       'host_is_superhost', 'host_thumbnail_url', 'host_picture_url',
       'host_neighbourhood', 'host_listings_count',
       'host_total_listings_count', 'host_verifications',
       'host_has_profile_pic', 'host_identity_verified', 'neighbourhood',
       'neighbourhood_cleansed', 'neighbourhood_group_cleansed', 'latitude'
       ,
       'longitude', 'property_type', 'room_type', 'accommodates', 'bathrooms',
       'bathrooms_text', 'bedrooms', 'beds', 'amenities', 'price',
       'minimum_nights', 'maximum_nights', 'minimum_minimum_nights',
       'maximum_minimum_nights', 'minimum_maximum_nights',
       'maximum_maximum_nights', 'minimum_nights_avg_ntm',
       'maximum_nights_avg_ntm', 'calendar_updated', 'has_availability',
       'availability_30', 'availability_60', 'availability_90',
       'availability_365', 'calendar_last_scraped', 'number_of_reviews',
       'number_of_reviews_ltm', 'number_of_reviews_l30d', 'first_review',
       'last_review', 'review_scores_rating', 'review_scores_accuracy',
       'review_scores_cleanliness', 'review_scores_checkin',
       'review_scores_communication', 'review_scores_location',
       'review_scores_value', 'license', 'instant_bookable',
       'calculated_host_listings_count',
       'calculated_host_listings_count_entire_homes',
       'calculated_host_listings_count_private_rooms',
       'calculated_host_listings_count_shared_rooms', 'reviews_per_month'],
      dtype='object')
```

Las columnas seleccionadas son las siguientes:

- latitude
- longitude
- accommodates
- bathrooms_text
- bedrooms
- availability_365
- number_of_reviews
- review_scores_value
- price

```
In [7]: # Filtramos las columnas que se consideran subjetivamente relevantes
data = data[['latitude', 'longitude', 'accommodates', 'bathrooms_text', 'bedrooms', 'availability_365', 'number_of_reviews', 'review_scores_value']]

In [8]: # Renombramos
data = data.rename({'review_scores_value': 'reviews_value'}, axis=1)

In [9]: # Clean numeric fields
data['bathrooms'] = data.bathrooms_text.str.extract('(\d+)')
data['bathrooms'] = data['bathrooms'].apply(pd.to_numeric, errors = 'coerce')
data['price'] = data['price'].replace('[\$,]', '', regex = True)
data['price'] = data['price'].apply(pd.to_numeric, errors = 'coerce')
print("Dataset tiene {} filas y {} columnas.".format(*data.shape))

Dataset tiene 17608 filas y 10 columnas.

In [10]: print("Shape: ", data.shape)
print("Column Information:", data.info())

Shape: (17608, 10)
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17608 entries, 11547 to 45499210
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
---  --  
 0   latitude          17608 non-null   float64
 1   longitude         17608 non-null   float64
 2   accommodates      17608 non-null   int64  
 3   bathrooms_text    17600 non-null   object  
 4   bedrooms          17333 non-null   float64
 5   availability_365  17608 non-null   int64  
 6   number_of_reviews  17608 non-null   int64  
 7   reviews_value     10949 non-null   float64
 8   price              17608 non-null   float64
 9   bathrooms          17585 non-null   float64
dtypes: float64(6), int64(3), object(1)
memory usage: 1.5+ MB
Column Information: None
```

```
In [11]: # Revisamos los NaN por columna  
data.isnull().sum()
```

```
Out[11]: latitude      0  
longitude      0  
accommodates    0  
bathrooms_text  8  
bedrooms        275  
availability_365 0  
number_of_reviews 0  
reviews_value    6659  
price            0  
bathrooms        23  
dtype: int64
```

```
In [12]: del data['bathrooms_text']
```

```
In [13]: # Reemplazamos los NaN por 1 baño, 1 habitación y 0 reviews  
data.fillna({'bathrooms': 1}, inplace = True)  
data.fillna({'bedrooms': 1}, inplace = True)  
data.fillna({'reviews_value': 0}, inplace = True)  
  
# Hacemos double-check  
data.isnull().sum()
```

```
Out[13]: latitude      0  
longitude      0  
accommodates    0  
bedrooms        0  
availability_365 0  
number_of_reviews 0  
reviews_value    0  
price            0  
bathrooms        0  
dtype: int64
```

```
In [14]: # Reordeno para dejar price la última columna  
data = data[['latitude', 'longitude', 'accommodates', 'bathrooms', 'bedrooms',  
            'availability_365', 'number_of_reviews', 'reviews_value', 'price']]
```

Eliminando propiedades no disponibles todo el año

En AirBnB el hospedador puede configurar un calendario de disponibilidad, de tal manera que la propiedad estará disponible unos días o semanas al año. Otras propiedades se encuentran disponibles todo el año, excepto cuando están reservadas. Para una mayor consistencia de los datos se aplicará la definición de "alta disponibilidad" establecida por AirBnB y que se establece en >60 días al año. Dicho lo cual se eliminarán las propiedades con disponibilidad <60 días y también a los nuevos alquileres con disponibilidad >300 días, para así centrar la aproximación.

La curva muestra una distribución bi-modal en disponibilidad, demostrando dos agrupaciones en los tipos de alquileres. Los picos de disponibilidad por debajo de 50 días indican alquileres de corto plazo, mientras que en el pico sobre los 365 días indican que son propiedades dedicadas íntegramente al alquiler. Una vez se han limpiado los datos, eliminare "availability_365" por su irrelevancia, tan solo revela los alquileres del próximo año.

In [15]:

```
# Distribución de availability_365
fig, axs = plt.subplots(ncols=2, figsize=(24, 4))
fig.suptitle('Distribución de disponibilidad (antes y después de eliminar los alquileres por temporadas)')


# Antes de la limpieza
x_axis = data['availability_365'].dropna()
sns.histplot(pd.Series(x_axis, name='Disponibilidad (antes)'), kde=True, ax=axs[0])
axs[0].set_title('Disponibilidad (antes)')

# Eliminar 60 < disponibilidad < 300 días
data = data.query('60 <= availability_365 <= 300')
print("Dataset tiene {} filas y {} columnas.".format(*data.shape))

# Despues de la limpieza
x_axis = data['availability_365'].dropna()
sns.histplot(pd.Series(x_axis, name='Disponibilidad (después)'), kde=True, ax=axs[1])
data = data.drop('availability_365', axis=1)
print("Dataset tiene {} filas y {} columnas.".format(*data.shape))
```

Dataset tiene 7343 filas y 9 columnas.

Dataset tiene 7343 filas y 8 columnas.



Eliminando grandes propiedades

Para este modelo eliminaré las propiedades con capacidad determinada por >10 huéspedes.

In [16]:

```
# Distribución de huéspedes
fig, axs = plt.subplots(ncols=2, figsize=(16, 4))
fig.suptitle('Distribución max invitados (antes y después de eliminar grandes alquileres)', weight='bold', fontsize=12)

# Antes de la limpieza
x_axis = data['accommodates'].dropna()
sns.histplot(pd.Series(x_axis, name='Max invitados (antes)'), kde=True, ax=axs[0])


# Eliminar huéspedes > 10
condition = data[data['accommodates'] > 10]
rows_to_drop = condition.index
print("{} filas eliminadas.".format(condition.shape[0]))
data = data.drop(rows_to_drop, axis=0)
print("Dataset tiene {} filas y {} columnas.".format(*data.shape))

# Despues de la limpieza
x_axis = data['accommodates'].dropna()
sns.histplot(pd.Series(x_axis, name='Max invitados (después)'), kde=True, ax=axs[1])
```

364 filas eliminadas.

Dataset tiene 6979 filas y 8 columnas.

Out[16]: <AxesSubplot: xlabel='Max invitados (después)', ylabel='Count'>



Eliminando las propiedades más caras

La gráfica izquierda muestra una distribución right-skewed con un long tail, dado por las propiedades de mayor importe. Para este modelo eliminaré los alquileres por encima de los 400€/noche para mantener su comparabilidad.

In [17]:

```
# Distribución del precio
fig, axs = plt.subplots(ncols = 2, figsize = (16, 4))
fig.suptitle('Distribución del precio (antes y después de eliminar los outliers)')

# Antes de la limpieza
x_axis = data['price'].dropna()
sns.histplot(pd.Series(x_axis, name = 'Precio (antes)'), kde = True, ax = axs[0])


# Eliminar precios > 400€/noche
condition = data[data['price'] > 400]
rows_to_drop = condition.index
print("{} filas eliminadas.".format(condition.shape[0]))
data = data.drop(rows_to_drop, axis = 0)
print("Dataset tiene {} filas y {} columnas.".format(*data.shape))

# Despues de la limpieza
x_axis = data['price'].dropna()
sns.histplot(pd.Series(x_axis, name = 'Precio (después)'), kde = True, ax = axs[1])
```

631 filas eliminadas.

Dataset tiene 6348 filas y 8 columnas.

Out[17]: <AxesSubplot: xlabel='Precio (después)', ylabel='Count'>



Estudio estadístico sobre precio

In [18]:

```
# Media
average_price = round(np.mean(data['price']), 3)

# Mediana
median_price = np.median(data['price'])

# Moda
mode_price = statistics.mode(data['price'])

# Varianza
variance_price = round(np.var(data['price']), 3)

# Desviación estandar
stdev_price = round(np.std(data['price']), 3)

# Primer cuartil
q1 = np.percentile(data.price, 25)


# Tercer cuartil
q3 = np.percentile(data.price, 75)

print('Media:\t\t', average_price)
print('Mediana:\t\t', median_price)
print('Moda:\t\t', mode_price)
print('Varianza:\t\t', variance_price)
print('Desviación estandar:\t', stdev_price)
print('Primer cuartil:\t', q1)
print('Tercer cuartil:\t', q3)

# Plot histogram
plt.hist(data.price, range = (0, 400), bins = 20, edgecolor = 'black')
plt.title("Análisis sobre el precio")
plt.xlabel('Precio (€)')
plt.ylabel('Disponibilidad en AirBnB')
plt.axvline(average_price, color = 'r', linewidth=1.5, label = "Media")
plt.axvline(median_price, color = 'y', linewidth=1.5, label = "Mediana")
plt.axvline(mode_price, color = 'orange', linewidth=1.5, label = "Moda")
plt.axvline(q1, color = '#6400e4', linewidth=1.5, label="Q1")
plt.axvline(q3, color = '#4fe0b0', linewidth=1.5, label="Q3")
plt.legend()

plt.show()
```

Media: 182.286
 Mediana: 175.0
 Moda: 150.0
 Varianza: 7861.528
 Desviación estandar: 88.665
 Primer cuartil: 115.0
 Tercer cuartil: 242.355



Representación sobre plano

- Mapa de calor: Representa zonas calientes por densidad de oferta.
- Clustering: Densidad de ofertas publicadas.
- Código de colores: Representando diferentes precios.

In [19]:

```
import folium
from folium.plugins import MarkerCluster, HeatMap


init_lat = data['latitude'].mean()
init_long = data['longitude'].mean()

m = folium.Map(
    tiles='CartoDB Dark_Matter',
    location=[init_lat, init_long],
    zoom_start=13,
    max_bounds=True
)
marker_cluster = MarkerCluster().add_to(m)
for index, row in data.iterrows():
    lat_air = row.latitude
    long_air = row.longitude
    price = row.price
    if (price > 0) & (price < 30):
        folium.Marker([lat_air, long_air],
                      tooltip=price,
                      icon=folium.Icon(color='green'),
                      clustered_marker = True
                     ).add_to(marker_cluster)
    elif (price < 50) & (price >= 30):
        folium.Marker([lat_air, long_air],
                      clustered_marker = True,
                      tooltip=price,
                      icon=folium.Icon(color='orange')
                     ).add_to(marker_cluster)
    elif (price < 90) & (price >= 50):
        folium.Marker([lat_air, long_air],
                      clustered_marker = True,
                      tooltip=price,
                      icon=folium.Icon(color='lightred')
                     ).add_to(marker_cluster)
    else:
        folium.Marker([lat_air, long_air],
                      clustered_marker = True,
                      tooltip=price,
                      icon=folium.Icon(color='red')
                     ).add_to(marker_cluster)

# convert to (n, 2) nd-array format for heatmap
privateArr = data[['latitude', 'longitude']].values

# plot heatmap
m.add_child(HeatMap(privateArr, radius=12))
```

Out[19]: Make this Notebook Trusted to load map: File -> Trust Notebook




In [20]:

```
# Crea matriz de correlación
_, ax = plt.subplots(figsize=(8, 8))

# Ploteo con Seaborn
sns.heatmap(data.corr(), cmap='coolwarm', center=0, square=True, cbar_kws={})
```

Out[20]: <AxesSubplot:>



In [21]:


```
# Filtra aquellos valores que tengan un 0.5 > abs > 1, los fusiona y elimina
highly_corr = data.corr()[(np.abs(data.corr()) > 0.5) & (np.abs(data.corr()) < 1.0)].drop('price', axis=1)

# Saco en pantalla los valores únicos
print(highly_corr['variable'].unique())

# Ploteo con Seaborn
sns.pairplot(data[highly_corr['variable'].unique()])
```

```
[ 'accommodates' 'bathrooms' 'bedrooms' 'price' ]
```

```
Out[21]: <seaborn.axisgrid.PairGrid at 0x7ffe7eb25340>
```



```
In [22]:
```

```
# Datos finales con los que trabajaremos el modelo  
data
```

Out[22]:

	latitude	longitude	accommodates	bathrooms	bedrooms	number_of_reviews
id						
159218	39.73689	2.89745	3	1.0	1.0	268
160482	39.65414	2.62546	6	2.0	3.0	62
166820	39.67932	2.50136	2	1.0	1.0	159
193426	39.89730	3.07411	6	2.0	3.0	82
210156	39.83375	3.10918	6	2.0	3.0	92
...
45459698	39.84032	3.12234	5	1.0	3.0	0
45475973	39.55374	2.61833	2	1.0	1.0	0
45478905	39.58170	2.67153	2	1.0	1.0	0
45493152	39.75437	2.90504	6	2.0	3.0	0
45496032	39.54550	2.39348	2	1.0	1.0	0

6348 rows × 8 columns

Preparación del modelo

In [23]:

```
# Preparo los datos para el modelo
Xraw = data.iloc[:, 0:7]
y = data.iloc[:, 7]
```

In [24]:

```
# Función para el cálculo del error
def RMSE(y_true, y_pred):
    return np.sqrt(mean_squared_error(y_true, y_pred))
```

Random Forest

Planteo un modelo de regresión basado en el algoritmo Random Forest, bajo el criterio de que tenemos muchas características a ser evaluadas y este modelo es bueno trabajando con variables independientes.

Selección de filtros

Dado que el número de variables a analizar es considerable, podemos intentar ajustarlas todas o buscar y ajustar las más relevantes. Para hacer esto último podemos seleccionar las mejores fuentes basadas en test estadísticos, eliminando todo excepto las de mayor resultado k .

In [25]:

```
results = []
count = 0
minim_abs = 99999999999
nmax = 9

N = Xraw.shape[1]
print(' |-----|-----| ')
print(' | k | RMSE | ')
for i in range (1,N):
    X = Xraw.copy()
    X = SelectKBest(f_regression,k=i).fit_transform(X,y)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
    random_forest_regression = RandomForestRegressor()
    random_forest_regression.fit(X_train,y_train)
    y_test_pred = random_forest_regression.predict(X_test)
    results.append((mean_squared_error(y_test,y_test_pred))**(1/2))
    print(' |-----|-----| ')
    print(' |',i,' | ',(mean_squared_error(y_test,y_test_pred))**(1/2), '| ')
    #Exit at nmax items not decreasing
    if len(results) != 0:
        if (mean_squared_error(y_test,y_test_pred))**(1/2) > minim_abs:
            count += 1
        else:
            minim_abs = (mean_squared_error(y_test,y_test_pred))**(1/2)
            count = 0
    else:
        count = 0
    if count > nmax:
        print(' |-----|-----| ')
        print(' | breaking after nmax items| ')
        break
print(' |-----|-----| ')

error_RF = RMSE(y_test, y_test_pred)
```

k	RMSE
1	68.30887410266877
2	64.95680837685326
3	64.69709491762613
4	66.03296757253719
5	66.455409681193
6	65.5799570786101

In [26]:

```
k = results.index(min(results)) + 1
print('El mínimo elemento en la lista es', k, 'y es', min(results))
```

El mínimo elemento en la lista es 3 y es 64.69709491762613

Model pre-filtered and score

In [27]:

```
N_ESTIM = 1000
X = Xraw
selection = SelectKBest(f_regression,k=k)
X = selection.fit_transform(X,y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)
random_forest_regression = RandomForestRegressor(n_jobs=2,n_estimators=N_ESTIM)
random_forest_regression.fit(X_train,y_train)
y_test_predict = random_forest_regression.predict(X_test)
y_train_predict = random_forest_regression.predict(X_train)
print('Valor de K: ', k)
print('Número de Random Forests: ', N_ESTIM)
rf_results=pd.DataFrame({'algorithm':['Random Forest'],
                         'Training error':[mean_absolute_error(y_train,y_train_predict)],
                         'Test error':[mean_absolute_error(y_test,y_test_predict)],
                         'Train_r2_score':[r2_score(y_train,y_train_predict)],
                         'Test_r2_Score':[r2_score(y_test,y_test_predict)]})
print(rf_results)
```

```
Valor de K:  3
Número de Random Forests:  1000
      algorithm  Training error  Test error  Train_r2_score  Test_r2_Score
0  Random Forest        50.190713    51.066497      0.477261      0.484015
```

GridSearch para RandomForest

In [28]:

```
#Grid model object
rf_grid = RandomForestRegressor()

param_grid = {'n_estimators':[1000],
              'max_leaf_nodes':[3,6,9],
              'ccp_alpha':[0.9,0.8,0.7]}

optimal_params = GridSearchCV(estimator = rf_grid,
                               param_grid=param_grid,
                               verbose=2,
                               n_jobs = 3,
                               cv = 3)

optimal_params.fit(X_train, y_train)

print(optimal_params.best_score_)
print(optimal_params.best_params_)
```

```
Fitting 3 folds for each of 9 candidates, totalling 27 fits
[Parallel(n_jobs=3)]: Using backend LokyBackend with 3 concurrent workers.
[Parallel(n_jobs=3)]: Done 27 out of 27 | elapsed: 41.7s finished
0.4536852609588826
{'ccp_alpha': 0.7, 'max_leaf_nodes': 9, 'n_estimators': 1000}
```

Decision Tree

```
In [29]: # Decision Tree
X = Xraw
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
regtree = DecisionTreeRegressor(min_samples_split=2, min_samples_leaf=4, max_depth=10)
dt = regtree.fit(X_train,y_train)
train_predict=dt.predict(X_train)
test_predict=dt.predict(X_test)
dt_results=pd.DataFrame({'algorithm':['Decision Tree'],
                           'Training error':[mean_absolute_error(y_train,train_predict)],
                           'Test error':[mean_absolute_error(y_test,test_predict)],
                           'Train_r2_score':[r2_score(y_train,train_predict)],
                           'Test_r2_Score':[r2_score(y_test,test_predict)]})
print(dt_results)


error_DeTr = RMSE(y_test, test_predict)
```

	algorithm	Training error	Test error	Train_r2_score	Test_r2_Score
0	Decision Tree	43.758183	49.451333	0.584901	0.500202

Ejemplo de representación de train y test entre las variables accommodates y price

```
In [30]: plt.scatter(X_train['accommodates'], y_train, label='train', c='firebrick', alpha=0.5)
plt.scatter(X_test['accommodates'], y_test, label='test', alpha=0.5)
plt.legend()
```

```
Out[30]: <matplotlib.legend.Legend at 0x7ffe7f603fd0>
```




Al aplicar decision tree, se puede jugar con tanto con su profundidad como con su número de hojas. Comparamos gráficamente como evoluciona el RMSE en función de la profundidad del arbol.

```
In [31]: max_depths = range(1, 20)
in_sample_errors = []
out_sample_errors = []
for max_depth in max_depths:
    tree = DecisionTreeRegressor(max_depth=max_depth).fit(X_train, y_train)
    y_pred_train = tree.predict(X_train)
    y_pred_test = tree.predict(X_test)
    in_sample_errors.append(RMSE(y_train, y_pred_train))
    out_sample_errors.append(RMSE(y_test, y_pred_test))

plt.plot(max_depths, in_sample_errors, c='firebrick', label='In-Sample Error')
plt.plot(max_depths, out_sample_errors, label='Out-Sample Error')

plt.xlabel('max_depth')
plt.ylabel('RMSE')
plt.legend(loc='best')
```

Out[31]: <matplotlib.legend.Legend at 0x7ffe7f42d0a0>



XGBoost

In [32]: xgb_round1 = XGBRegressor()

In [33]: param_grid = {
 'learn_rate':[0.1,0.05,0.01],
 'min_child_weight':[3,4,5],
 'gamma':[i/10.0 for i in range(3,6)],
 'subsample':[i/10.0 for i in range(6,11)],
 'colsample_bytree':[i/10.0 for i in range(6,11)],
 'max_depth': [9,10,11],
 'n_estimators':[300,500,1000]
}

In [34]:

```
optimal_params = GridSearchCV(  
    estimator = xgb_round1,  
    param_grid=param_grid,  
    verbose=2, # NOTE: If you want to see what Grid Search is doing  
    n_jobs = 2,  
    cv = 3  
)  
  
optimal_params.fit(X_train, y_train, early_stopping_rounds=10, eval_set=[(X_val, y_val)])  
  
print(optimal_params.best_score_)  
print(optimal_params.best_params_)
```

Fitting 3 folds for each of 6075 candidates, totalling 18225 fits

```
[Parallel(n_jobs=2)]: Using backend LokyBackend with 2 concurrent workers.  
[Parallel(n_jobs=2)]: Done  37 tasks      | elapsed:    8.9s  
[Parallel(n_jobs=2)]: Done 158 tasks      | elapsed:   28.0s  
[Parallel(n_jobs=2)]: Done 361 tasks      | elapsed:  1.0min  
[Parallel(n_jobs=2)]: Done 644 tasks      | elapsed:  1.9min  
[Parallel(n_jobs=2)]: Done 1009 tasks     | elapsed:  2.9min  
[Parallel(n_jobs=2)]: Done 1454 tasks     | elapsed:  4.1min  
[Parallel(n_jobs=2)]: Done 1981 tasks     | elapsed:  5.6min  
[Parallel(n_jobs=2)]: Done 2588 tasks     | elapsed:  7.3min  
[Parallel(n_jobs=2)]: Done 3277 tasks     | elapsed:  9.1min  
[Parallel(n_jobs=2)]: Done 4046 tasks     | elapsed: 11.1min  
[Parallel(n_jobs=2)]: Done 4897 tasks     | elapsed: 13.4min  
[Parallel(n_jobs=2)]: Done 5828 tasks     | elapsed: 15.8min  
[Parallel(n_jobs=2)]: Done 6841 tasks     | elapsed: 18.8min  
[Parallel(n_jobs=2)]: Done 7934 tasks     | elapsed: 22.0min  
[Parallel(n_jobs=2)]: Done 9109 tasks     | elapsed: 25.3min  
[Parallel(n_jobs=2)]: Done 10364 tasks    | elapsed: 28.9min  
[Parallel(n_jobs=2)]: Done 11701 tasks    | elapsed: 33.0min  
[Parallel(n_jobs=2)]: Done 13118 tasks    | elapsed: 37.4min  
[Parallel(n_jobs=2)]: Done 14617 tasks    | elapsed: 42.0min  
[Parallel(n_jobs=2)]: Done 16196 tasks    | elapsed: 47.1min  
[Parallel(n_jobs=2)]: Done 17857 tasks    | elapsed: 52.4min  
[Parallel(n_jobs=2)]: Done 18225 out of 18225 | elapsed: 53.5min finished
```

```
[19:49:04] WARNING: /Users/travis/build/dmlc/xgboost/src/learner.cc:541:  
Parameters: { learn_rate } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[0] validation_0-rmse:151.62137  
[1] validation_0-rmse:115.65456  
[2] validation_0-rmse:92.50348  
[3] validation_0-rmse:78.45510  
[4] validation_0-rmse:70.22861  
[5] validation_0-rmse:65.54177  
[6] validation_0-rmse:62.93241  
[7] validation_0-rmse:61.63047  
[8] validation_0-rmse:60.77943  
[9] validation_0-rmse:60.30464  
[10] validation_0-rmse:60.11770  
[11] validation_0-rmse:60.00397  
[12] validation_0-rmse:59.98875  
[13] validation_0-rmse:60.02403  
[14] validation_0-rmse:60.03123  
[15] validation_0-rmse:60.02966  
[16] validation_0-rmse:60.13486  
[17] validation_0-rmse:60.19898  
[18] validation_0-rmse:60.25098  
[19] validation_0-rmse:60.16614  
[20] validation_0-rmse:60.22897  
[21] validation_0-rmse:60.21773  
0.5233638198362006  
{'colsample_bytree': 1.0, 'gamma': 0.3, 'learn_rate': 0.1, 'max_depth': 9, 'min_child_weight': 5, 'n_estimators': 300, 'subsample': 1.0}
```

Final XGBoost model after parameters

```
In [35]:  
# Final model after adjusting parameters.  
xg = XGBRegressor(seed=42,  
                   objective='reg:squarederror',  
                   gamma=0.3,  
                   learn_rate=0.1,  
                   max_depth=9,  
                   min_child_weight=5,  
                   subsample=1,  
                   colsample_bytree=1,  
                   n_estimators=300)  
  
xg.fit(X_train, y_train,  
       verbose=False,  
       early_stopping_rounds=10,  
       eval_set=[(X_test, y_test)])
```

```
[19:49:05] WARNING: /Users/travis/build/dmlc/xgboost/src/learner.cc:541:  
Parameters: { learn_rate } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
Out[35]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                      colsample_bynode=1, colsample_bytree=0.6, gamma=0.3, gpu_id=-1  
,  
                      importance_type='gain', interaction_constraints='', learn_rate  
=0.1,  
                      learning_rate=0.300000012, max_delta_step=0, max_depth=9,  
                      min_child_weight=5, missing=nan, monotone_constraints='()',  
                      n_estimators=300, n_jobs=4, num_parallel_tree=1, random_state=42,  
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=42,  
                      subsample=1, tree_method='exact', validate_parameters=1,  
                      verbosity=None)
```

```
In [36]:  
# XGB Regressor  
train_predict=xg.predict(X_train)  
test_predict=xg.predict(X_test)  
xg_results=pd.DataFrame({'algorithm':['XGBoost'],  
                         'Training error':[mean_absolute_error(y_train, train_predict)],  
                         'Test error':[mean_absolute_error(y_test, test_predict)],  
                         'Train_r2_score':[r2_score(y_train, train_predict)],  
                         'Test_r2_Score':[r2_score(y_test, test_predict)]})  
print(xg_results)  
  
error_XGB = RMSE(y_test, test_predict)
```

	algorithm	Training error	Test error	Train_r2_score	Test_r2_Score
0	XGBoost	30.945974	46.334156	0.782058	0.557431

```
In [37]: pd.concat([rf_results, dt_results, xg_results], axis=0, ignore_index=True)
```

```
Out[37]:  
algorithm  Training error  Test error  Train_r2_score  Test_r2_Score  
0  Random Forest  50.190713  51.066497  0.477261  0.484015  
1  Decision Tree  43.758183  49.451333  0.584901  0.500202  
2  XGBoost  30.945974  46.334156  0.782058  0.557431
```

```
In [38]:  
print('Los errores con los distintos tipos de modelos optimizados son los siguientes: ')  
print('Error del modelo RF: {} k$'.format(e=round(error_RF, 2)))  
print('Error medio RF: {} k$'.format(e=round((error_RF/len(X_test))*1000, 2)))  
print('\nError del modelo DT: {} k$'.format(e=round(error_DeTr, 2)))  
print('Error medio DT: {} k$'.format(e=round((error_DeTr/len(X_test))*1000, 2)))  
print('\nError del modelo XGBoost: {} k$'.format(e=round(error_XGB, 2)))  
print('Error medio XGBoost: {} k$'.format(e=round((error_XGB/len(X_test))*1000, 2)))
```

Los errores con los distintos tipos de modelos optimizados son los siguientes:

Error del modelo RF: 65.58 k\$

Error medio RF: 51.64 \$

Error del modelo DT: 62.91 k\$

Error medio DT: 49.53 \$

Error del modelo XGBoost: 59.2 k\$

Error medio XGBoost: 46.61 \$

Conclusiones

El mejor modelo encontrado haciendo uso de diferentes técnicas es el XGBoost, con un error medio de 46,61\$.

Es realmente complejo estimar el valor de una propiedad basándonos en determinadas variables, por ejemplo, fueron descartados los ammenities debido a que incluían elementos tan dispares como piscina y tostador o parking y secador del pelo, con lo que la manera más exhaustiva sería asignar un valor a cada ammenitie, pero complicaría en exceso esta tarea.

Es interesante cómo al implementar los algoritmos y al optimizar sus hiperparámetros se ha disminuido el tiempo de procesamiento de 81 minutos a unos 53 minutos.